
The Essence of Reactive Programming

A Theoretical Approach

By

EDDY BERTOLUZZO



Faculty of Electrical Engineering, Mathematics and Computer Science
DELFT UNIVERSITY OF TECHNOLOGY

A dissertation submitted to Delft University of Technology in
accordance with the requirements of the MASTER OF SCIENCE
degree in the Faculty of Electrical Engineering, Mathematics and
Computer Science.

OCTOBER 2016

ABSTRACT

In the last few years the *** of Reactive Programming has gained much traction in the developer's community, especially with the general trend for programming languages to embrace functional programming and the shifting of today's applications towards an asynchronous approach to modeling data. The goal of this work is to shed some light to the sea of more or less informed opinions regarding what the meaning of Reactive Programming is thorough the use of Mathematics...

DEDICATION AND ACKNOWLEDGEMENTS

*"Well, here at last, dear friends,
on the shores of the Sea comes the
end of our fellowship in
Middle-earth. Go in peace! I will not
say: do not weep, for not all tears are
an evil."*

— J.R.R. Tolkien,
The Return of the King

H ere goes the dedication.

TABLE OF CONTENTS

	Page
1 Reactive Programming	5
1.1 The Essence of Reactive Programs	5
1.2 Why Reactive Programming Matters	6
1.3 Reactive Programming in the Real World	7
1.3.1 Reactive Extensions	7
1.3.2 Reactive Streams	7
1.3.3 Functional Reactive Programming	8
2 Into the Rabbit Hole: <i>Deriving the Observable</i>	9
2.1 Iterables	10
2.2 The Essence of Iterables	11
2.3 Applying Duality	14
2.4 Termination and Error Handling	17
2.5 Observables are Continuations	18
3 Out of the rabbit hole: <i>Towards a usable API</i>	23
3.1 The Reactive Contract	24
3.2 Concurrency with Schedulers	24
3.2.1 A note on the concept of time	28
3.2.2 Impact on Operators	28
3.3 Subscriptions	28
3.4 Operators	33
A Appendix A	39
A.1 Categorical Duality	39
A.2 Continuations	39
A.3 (Co)Products	39
A.4 (Un)Currying	39
A.5 (Covariance & Contravariance)	39
A.6 Functors	39

TABLE OF CONTENTS

B Appendix B	41
Bibliography	43

INTRODUCTION

*"Lasciate ogni speranza, voi
ch'intrate."*

— Dante Alighieri,
Divina Commedia

With the evolution of technologies brought in by the new millennium and the exponential growth of Internet-based services targeting millions of users all over the world, the Software Engineers' community has been continuously tested by an ever growing number of challenges related to management of increasingly large amounts of user data^[6].

This phenomena is commonly referred to as Big Data. A very popular 2001 research report^[10] by analyst Doug Laney, proposes a definition of big data based on its three defining characteristics:

- *Volume*: the quantity of data applications have to deal with, ranging from small - e.g. locally maintained Databases - to large - e.g. distributed File Systems replicated among data centers.
- *Variety*: the type and structure of data, ranging from classic SQL-structured data sets to more diversified and unstructured ones such as text, images, audio and video.
- *Velocity*: the speed at which data is generated, establishing the difference between pull-based systems, where data is synchronously pulled by the consumer, and push-based systems, more suited for handling real-time data by asynchronously pushing it to its clients.

Each of these traits directly influences the way programming languages, APIs and databases are designed today. The increasing volume calls for a declarative approach to data handling as opposed to an imperative one, resulting in the developer's focus shifting from how to compute something to what it is to be computed in the first place^[5]. The diversification of data, on the other hand, is the main drive for the research and development of noSQL approaches to data storage. Lastly, the increase in velocity fuels the need for event-driven, push-based models of computation that can better manage the high throughput of incoming data^[11].

In this context, the concept of *reactive programming* has gained much traction in the developer's community as a paradigm well-suited for the development of asynchronous event-driven applications^[2]. Unfortunately, reactive programming has been at the center of much discussion, if not confusion, with regards to its definition, properties and principles that identify it^[12].

The goal of our work is to use mathematics as a tool to analyze today's commercial reactive libraries and understand how they relate to the theoretical concept of reactive programming. We are going to do so by utilizing concepts and ideas from functional programming and category theory with the purpose of deriving the reactive types, and will subsequently continue with the development of a reference reactive library which builds upon the previously derived theoretical foundations.

Contributions

To the best of our knowledge, we are not aware of any previous work which analyses reactive programming from a theoretical point of view or derives its types and interfaces through the use of mathematics.

The most known attempt at defining reactive programming is the Reactive Manifesto^[7], a document that tries to describe reactive systems in terms of design principles and conceptual architecture. Although certainly insightful, we feel the document targets a more managing-focused audience as opposed to software developers and engineers.

Our work, on the other hand, aims at providing types and interfaces that describe the real essence of the reactive paradigm, aiding engineers that wish to use or develop reactive libraries in understanding and taking more informed decisions in the matter.

Overview

Chapter 2 starts with ... bla bla bla

Write overview

Notation & Conventions

In the exposition of our work we will make use of Haskell as the reference language. This decision is motivated by the language's strong connection with mathematics as well as its clean syntax.

All the code presented in this report, a minimal complete theoretical implementation and a full blown library implementation of a reactive library can be found at the associated code repository on Github ().

Github link

REACTIVE PROGRAMMING

The cold winds are rising in the North... Brace yourselves, winter is coming.

— George R.R. Martin,
A Game of Thrones

In this chapter we are going to introduce the concept of reactive programming and motivate its importance and relevance with regards to modern applications and the type of problems developers have to face nowadays. We are then going to introduce the most popular commercial libraries that claim to solve the reactive problem, with the purpose of giving the reader some context for our discussion and motivating the need for a mathematical formalization that abstracts over the class of problems these implementations set out to address.

1.1 The Essence of Reactive Programs

The use of the term reactive program in scientific literature is dated back to the mid-sixties^[1]. A relevant and insightful definition was given by G. Berry in 1991^[3] as he describes reactive programs in relation to their dual counterparts, interactive programs:

“Interactive programs interact at their own speed with users or with other programs; from a user point of view, a time-sharing system is interactive. Reactive programs also

maintain a continuous interaction with their environment, but at a speed which is determined by the environment, not by the program itself.”

Interactive programs concretize the idea of a pull-based model of computation, where the program - the consumer in this case - has control over the speed at which data will be requested and handled. A perfect example of an interactive program is a control-flow structure such as a for-loop iterating over a collection of data: the program is in control of the speed at which data is retrieved from the containing collection and will request the next element only after it is done handling the current one.

Reactive programs, on the contrary, embody the idea of a push-based - or event-driven - model of computation, where the speed at which the program interacts with the environment is determined by the environment rather than the program itself. In other words, it is now the producer of the data - i.e. the environment - who determines the speed at which events will occur whilst the program's role reduces to that of a silent observer that will react upon receiving events. Standard example of such systems are GUI applications dealing with various events originating from user input - e.g. mouse clicks, keyboard button presses - and programs dealing with stock market, social media or any other kind of asynchronous updates.

1.2 Why Reactive Programming Matters

Considering the definition and examples of reactive programs we analyzed in the previous section, let's now try to formalize the class of problems the reactive programming paradigm is specifically well-suited for.

The table below provides a collection of types offered by common programming languages for handling data, parameterized over two variables: the size of the data, either one or multiple values, and the way data is handled, either by synchronous or asynchronous computations^[13].

	One	Many
Sync	<code>a</code>	<code>Iterable a</code>
Async	<code>Future a</code>	<i>Reactive Programming</i>

The first row shows that synchronous functions come in two flavors: classic functions that return a single value of type `a` and functions that produce a collection of results of type `a`, abstracted through the `Iterable a` interface (See section 2.1). These types of functions embody the standard imperative, pull-based approach to programming, where a call to a function/method synchronously blocks until a result is produced.

Moving on to the second row, we encounter `Future a`, an interface representing an asynchronous computation that, at a certain point in the future, will result in a value of type `a`. Futures are generally created by supplying two callbacks together with the asynchronous computation, one to be executed in case of success and the other one in case of error.

Programming languages, however, are not as well equipped when it comes to handling asynchronous computations resulting in multiple values - i.e. push-based collections. The issue lies in the fact that the program's control flow is dictated by the environment rather than the program itself - i.e. inversion of control -, making it very hard to model such problems with commonly known control structures, which are optimized for sequential models of computation. Traditional solutions typically involve developers manually trying to compose callbacks by explicitly writing CPS (continuation passing style) code^[13], resulting in what it's commonly referred to as *Callback Hell*^[4].

The aforementioned class of problems reflects the definition of reactive programs we analyzed in the previous section, where the environment asynchronously - i.e. at its own speed - pushes multiple events to the program. The reactive programming paradigm sets out to provide interfaces and abstractions to facilitate the modeling of such problems as push-based collections.

However, interfaces are only as good as the implementations that back them up. In the next section we are going to discuss and analyze various APIs and libraries that claim to embody the reactive paradigm, motivating our need for a mathematical formalization to aid in unifying these different approaches under a single set of interfaces, eliminating doubts and confusion when it comes to the essence of reactive programming.

1.3 Reactive Programming in the Real World

Talk about FRP, Reactive Streams, Rx and other stuff. Cite survey paper Claim that Observable is dual of Iterator in order to justify the why.

1.3.1 Reactive Extensions

1.3.2 Reactive Streams

Try to define itself as reactive programming, mathematics prove the claim wrong, they implement asynchronous iterables. This does not make them useless, yet they are not a good solution for the problem we are trying to solve.

Probably put the derivation of Reactive Streams in the appendix, it's not really the focus, I guess.

1.3.3 Functional Reactive Programming

Much researched topic, beautiful in theory, yet based on continuous time, un-achievable with computers where time is inherently discrete. Connel Elliot, the inventor of FRP, says that commercial implementations are far from the original theory and do not really represent the essence of FRP

INTO THE RABBIT HOLE: *Deriving the Observable*

"It was much pleasanter at home," thought poor Alice, "when one wasn't always growing larger and smaller, and being ordered about by mice and rabbits. I almost wish I hadn't gone down the rabbit-hole – and yet – and yet – ..."

— Lewis Carol,
Alice in Wonderland

As we saw in Chapter 1, the `Iterable` interface embodies the idea of a pull-based model of computation and is the commonly adopted solution to dealing with synchronous computations resulting in multiple values. In this chapter we are going to formalize the intuition that there exists a duality relation between interactive and reactive programs, as well as between pull and push models of computations, by deriving the `Observable` interface starting from its dual counterpart, the `Iterable`.

The derivation that follows will require the use of a number mathematical concepts such as *categorical duality*, *continuations*, *(co)products*, *(un)currying*, *covariance*, *contravariance* and *functors*. We suggest the reader to get familiar with these topics before diving into the derivation. An accessible introduction to each can be found in Appendix A.

2.1 Iterables

An `Iterable` is a programming interface which enables the user to traverse a collection of data, abstracting over the underlying implementation^[7].

The interface and semantics of `Iterables` were first introduced by the Gang of Four through their Iterator pattern^[7]; today's most used programming languages introduce the `Iterable` as the root interface for standard pull collections APIs - exposing concrete implementations such as maps, sets, indexed sequences and so on.

The `Iterable` interface is generally fixed across programming languages, with the exception of naming conventions - C# and related languages call it `IEnumerable` - and slight differences in the types, as we can see from the following definitions.

```
1  -- Java Iterable
2  class Iterable a where
3      getIterator :: () -> IO (Iterator a)
4
5  class Iterator a where
6      hasNext :: () -> Bool
7      next    :: () -> IO a
8
9  -----
10
11 -- C# IEnumerable
12 class IEnumerable a where
13     GetEnumerator :: () -> IO (IEnumerator a)
14
15 class IEnumerator a where
16     moveNext :: () -> IO Bool
17     current  :: () -> a
```

Although the essence of the pattern is preserved by both definitions, we claim that the C# version more clearly and accurately reflects the way side effects play a role in the usage of the interface: `moveNext` contains all the side effects of traversing the underlying collection and retrieving the next value while `current` can inspect the retrieved value multiple times in a pure way. The Java version, on the other hand, embeds the side effect in the `next` function, making it impossible to inspect the current value multiple times. For this reason and without loss of generality, we will make use of the

C# definition - modulo naming conventions - in the reminder of the discussion:

```
1 class Iterable a where
2     getIterator :: () -> IO (Iterator a)
3
4 class Iterator a where
5     moveNext :: () -> IO Bool
6     current  :: () -> a
```

2.2 The Essence of Iterables

The first step in deriving the `Observable` is to simplify our `Iterable` definition to a type that reflects its very essence; we are gonna do this by stripping the interface presented in the previous section of all the unnecessary operational features that only clutter our definition.

Let's start by taking a closer look at the `Iterator` interface; we can observe that the definition of the functions `moveNext` and `current` is equivalent to a single function which returns either a value - analogous to a `moveNext` call returning true and a subsequent invocation to `current` - or nothing - analogous to a call to `moveNext` returning false.

Before we formalize this observation with a proper type, let us notice another effect that is hidden in the current definition of `moveNext` and not made explicit by its type: the possibility for an exception to be thrown by the function's body.

By merging these considerations with the notion of coproducts and Haskell's `Either` and `Maybe` type, we obtain the following definition.

```
1 class Iterable a where
2     getIterator :: () -> IO (Iterator a)
3
4 class Iterator a where
5     moveNext :: () -> IO (Either SomeException (Maybe a))
```

Note how, theoretically, `getIterator` could also throw an exception. We assume here, without loss of generality, that the function will never throw and will always be able to return an `Iterator` instance.

The next step is to forget about class instances and express our interfaces as simple types.

```
1 type Iterable a = () -> IO (Iterator a)
2 type Iterator a = () -> IO (Either Exception (Maybe a))
```

At this point, we want to put aside the operational concerns regarding exceptions and termination and assume the `Iterator` function will always return a value of type `a`.

```
4 type Iterable a = () -> IO (() -> IO a)
5 type Iterator a = () -> IO a
```

We have now reached a point where no simplification is possible anymore. The obtained types reflect the essence of the Iterator pattern: an `Iterable` is, theoretically, a function which, when invoked, produces an `Iterator` and an `Iterator` is itself a function producing a value of type `a` as a side effect.

When looking at the `Iterator` type from an object oriented perspective, the reader should notice a strict similarity to a *getter* function - i.e. a lazy producer of values: iterators are, in fact, nothing more than getters of values of type `a`. The `Iterable`, on the other hand, is a function that enables the user to get an `Iterator`, i.e. a getter of a getter of `a`.

This correspondence will turn out to be very insightful later on in our discussion, where we will observe that `Observable` is nothing more than a setter of setters, another instance of duality in our formalization.

When looking at the relation between the `Iterator` type and its base component, `a`, we can observe how they are bound by a covariance relation:

$$\frac{A <: B}{() \rightarrow A <: () \rightarrow B}$$

The intuition can be easily understood when we think of an iterable as a soft-drink vending machine, i.e. a function which, whenever called, will give back a soft-drink:

$$\frac{Coke <: SoftDrink}{() \rightarrow Coke <: () \rightarrow SoftDrink}$$

If coke is a subtype of soft-drink, then whenever I am asked for a soft-drink vending machine, I can

hand out a coke vending machine without incurring in any troubles with the person who asked, as that machine will correctly provide soft drinks whenever prompted for one.

With `Iterable` being a getter itself, it should be clear how covariance plays the same role as with `Iterator`.

To formally prove the intuition of a covariant relation, we instantiate the `Iterable` / `Iterator` types to a covariant `Functor`.

```

3 {-
4   Using a simplified syntax, hiding IO and Haskell's newtypes,
5   in order to clearly show the idea:
6
7   iterator :: () -> a
8   iterable :: () -> () -> a
9
10  fmapi :: (a -> b) -> (() -> a) -> () -> b
11  fmapi      f          ia      = () -> f (ia ())
12
13  fmapii :: (a -> b) -> (() -> () -> a) -> () -> () -> b
14  fmapii      f          iia     = () -> fmapi f (iia ())
15 -}
16
17 newtype Iterator a = Iterator { runIterator :: () -> IO a }
18 newtype Iterable a = Iterable { getIterator :: () -> IO (Iterator a) }
19
20 instance Functor Iterator where
21   fmap f ia = Iterator $ \_ -> liftM f (runIterator ia ())
22
23 instance Functor Iterable where
24   fmap f iia = Iterable $ \_ -> liftM (fmap f) (getIterator iia ())

```

We will see in the next section how these concepts are relevant in expressing and motivating the duality between Iterables and Observables.

2.3 Applying Duality

By now, the reader should be somehow familiar with the concept of duality, as it has appeared many times throughout our discussion in concepts such as pull and push models of computation or interactive and reactive programs. Duality is, in fact, a very important general theme that has manifestations in almost every area of mathematics^[8] (See Appendice A for an introductory discussion on the topic).

Starting from the fact that the `Iterable` interface embodies the idea of interactive programming, let's use the principle of duality to derive the `Observable` interface and see how it relates to the concept of reactive programming. In practice, this translates to the simple task of flipping the function arrows in the `Iterable` interface, taking us from a function resulting in a value of type `a` to one accepting an `a`.

```

1  {-
2      () -> (() -> a) -- iterable
3      () <- (() <- a) -- apply duality
4      (a -> ()) -> () -- observable
5  -}
6
7  type Iterator a = () -> IO a
8              -- = () IO <- a
9  type Observer a = a -> IO ()
10
11 type Iterable  a = () -> IO (() -> IO a)
12              -- = () IO <- (() IO <- a)
13 type Observable a = (a -> IO ()) -> IO ()

```

Note how the side effects are bound to function application rather than values, hence their flipped position in the `Observable` type.

The reader acquainted with functional programming will easily see the resemblance between the `Observable` type and a CPS function (See Appendice A).

```

1  cont      :: (a -> r) -> r
2  observable :: (a -> IO ()) -> IO ()

```

It is clear how an `Observable` is nothing more than a CPS function where the result type `r` is instantiated to `IO ()`. To convince ourselves of this equivalence, let's think about the definition of a CPS function, i.e. a suspended computation which, given another function - the continuation - as argument, will produce its final result. This definition suits perfectly the idea behind `Observable` discussed in Section 1.3.1: a function which will do nothing - i.e. is suspended - until it is subscribed to by an `Observer`.

A continuation, on the other hand, represents the future of the computation, a function from an intermediate result to the final result^[14]; in the context of `Observable`s, the continuation represents the `Observer`, a function specifying what will happen to a value produced by the `Observable`, whenever it will become available, that is, whenever it will be pushed into the `Observer`. Since a continuation can be called multiple times within the surrounding CPS context, it is easy to see how this mathematical concept allows us to deal with multiple values produced at different times in the future.

In the previous section we have discussed many properties associated with `Iterable`s. Let's analyze now how these properties translate under dualisation and how they affect our new derived interface, the `Observable`.

First, moving from the observation that an `Iterable` is a getter of a getter, we can observe that the `Observable` plays exactly the opposite role, that is, a setter of a setter. The type `Observer :: a -> IO ()` represents, in fact, the essence of a setter function, whereas the `Observable` consists in nothing more than the simple task of applying the observer function to itself, producing a setter of setters.

While the discussion about `Iterable`'s covariance was quite intuitive, things get a little bit more complicated when analyzing `Observable`s. Let's start by informally introducing the notion of positivity and negativity of types: we can interpret a function of type `f :: a -> b` as a way for us to produce a value of type `b`. In this context, `b` is considered to be positive with respect to the type `a -> b`. On the other hand, in order to apply the function, we are going to need a value of type `a`, which we will need to get from somewhere else; `a` is therefore considered to be negative w.r.t. the function type, as the function introduces a need for this value in order to produce a result. The point of this distinction is that positive type variables introduce a covariant relation between base and function type whereas negative type variables introduce a contravariant relation.

Analyzing `Iterable` within this framework is easy, the `Iterator` function contains a single type parameter found in a positive position, therefore resulting in a covariant relation; being the `Iterable` the result of applying the `Iterator` function to itself, we again result in a covariant relation w.r.t. the type parameter `a`.

The `Observer` function, on the contrary, introduces a need for a value of type `a`, resulting in a con-

travariant relation w.r.t. `a`. Again, the `Observable` function is the result of applying `Observer` to itself; surprisingly, this results in `a` being in a positive position. The intuition is easily understood by thinking about the rules of arithmetic multiplication: `a` is in negative position w.r.t. the `Observer` function, whereas the `Observer` is in negative position w.r.t. the `Observable`. This leads to `a` being negated twice, ultimately resulting in a positive position within the `Observable` function.

```
3  f  :: a -> b
4      = -a -> +b
5
6  g  :: (a -> b) -> c
7      = -(-a -> +b) -> +c
8      = (+a -> -b) -> +c
9
10 observer
11     :: a -> ()
12     = -a -> ()
13
14 observable
15     :: (a -> ()) -> ()
16     = -(-a -> ()) -> ()
17     = (+a -> ()) -> ()
```

Before we formalize this claim, let's convince ourselves that `Observable` s effectively produces a value of type `a` by looking at an example:

```
21 randomValueObs :: Observable Int
22 randomValueObs = Observable $ \observer -> do
23   int <- randomRIO (1, 10)
24   observer int
```

It is clear from this implementation that `randomValueObs` indeed produces a value of type `Int`, whereas the `Observer` introduces a need for such value in order to be applied. For more details on the positivity and negativity of functions and type variables, see ^[15] [9].

Just as we did with `Iterable/Iterator`, we can formally prove the covariant and contravariant re-

lations between `Observable/Observer` and their base type `a` by instantiating them to `Functor` and `Contravariant` - i.e. `Cofunctor` - respectively.

```

3 newtype Observer a = Observer { onNext :: a -> IO () }
4 newtype Observable a = Observable { subscribe :: Observer a -> IO () }
5
6 instance Contravariant Observer where
7     contramap f ob = Observer $ onNext ob . f
8
9 instance Functor Observable where
10     fmap f ooa = Observable $ subscribe ooa . contramap f
11

```

2.4 Termination and Error Handling

We began this chapter by progressively simplifying the `Iterable`'s interface in order to derive a type that would theoretically represent its very essence. One of the most important steps was setting aside concerns regarding termination and error handling of a collection. We are now going to reshape our reactive interfaces in order to address these concerns and appropriately describe the potential side effects directly in the types.

Informally, an `Observable` stream might not only produce one or more values, but it might gracefully terminate at a certain point in time or throw an exception and abruptly terminate whilst processing values. A more appropriate type for `Observer` is then the following:

```

1 newtype Observer a = Observer
2   { onNext :: Either SomeException (Maybe a) -> IO ()
3   }

```

Just as with `Iterable`, the introduction of `Either SomeException` allows us to express that the `Observer` can handle unexpected exceptions, while the `Maybe` reflects the possibility for a stream to end and propagate no more values.

Unfortunately, this type is very hard to read as well as understand for someone new to the topic. Looking at the matter from a functionality point of view, what we really need is for our CPS func-

tion - i.e. the `Observable` - to accept two additional continuations, one dealing with completion and one with exceptions. We can achieve this by first noticing that our type is nothing more than a coproduct - the same that we introduced previously for `Iterable` - of three base types: `a + SomeException + ()`. By utilizing the notion of product - the dual of coproduct - we can split the function handling the initial type into three different ones. This brings us to the final version of our reactive interfaces for push-base collections¹

```
1 newtype Observable a = Observable
2   { subscribe :: Observer a -> IO ()
3   }
4
5 data Observer a = Observer
6   { onNext      :: a -> IO ()
7   , onError     :: SomeException -> IO ()
8   , onCompleted :: IO ()
9   }
```

The `Observable` is now a special version of a CPS function accepting three continuation functions - embedded inside the `Observer` -, one for each effect an `Observable` can propagate: value, termination or exception.

2.5 Observables are Continuations

Throughout our discussion, we have mentioned multiple times the connection between `Observable` and CPS functions. In this section we are going to explore this matter and see how these two concepts are effectively the same thing.

We are going to start from the notion that `Observable` is, at its essence, nothing more than a setter of setters, the result of applying the `Observer` function to itself. We can then express the `Observer` as a function `(!)`² that negates its type argument and results in a side-effectful computation.

```
1 !a :: a -> IO ()
```

¹The two definitions are equivalent also from an implementation point of view, the first simulating the second through the use of pattern matching.

²Regard the code used in this explanation as pseudo-Haskell.

When we apply the function to itself - i.e. substitute `a` for `!a` - we obtain our first definition of `Observable`, a CPS function that instantiates the result to `IO ()`.

```
1 !!a :: (a -> IO ()) -> IO ()
```

As we have seen in the previous section, this definition is not expressive enough when we want to make explicit all the effects that are involved when dealing with push-base collections. It is therefore necessary to deviate from the standard definition of continuation and replace the inner application of `(!)` with a new function `(?)`, whose type embed the involved effects:

```
1 ?a :: Either Error (Maybe a) -> IO () -- termiantion and error handling
2 !?a :: (Either Error (Maybe a) -> IO ()) -> IO ()
```

Note how this definition is equivalent to the one used in the previous section, where we used the notion of products to unwrap the `(?)` function into three different ones, each addressing one of the possible effects.

Finally, we can now translate this pseudo-code into real Haskell, proving the equivalence between `Observable` and CPS functions. We are going to do so with the help of Haskell's continuation monad transformer, a construct which allows to stack the functionality of the continuation monad on top of other monads. With it, we can define the `Observable` as a continuation transformer producing an empty result in the IO monad.

```
5  --   Event a = Either SomeException (Maybe a)
6  data Event a = OnNext a | OnError SomeException | OnCompleted
7      deriving Show
8
9  type Observer a = Event a -> IO ()
10 type Observable a = ContT () IO (Event a)
11
12 newObservable :: (Observer a -> IO ()) -> Observable a
13 newObservable = ContT
14
15 subscribe :: Observable a -> Observer a -> IO ()
16 subscribe = runContT
```

The code above uses a slightly different approach to expressing the three types of side effects an `Observer` has to deal with. Instead of using `Either` and `Maybe` from Haskell's libraries, we utilize our own custom datatype `Event`, a coproduct of values of type `a + SomeException + ()`; although the two definitions are equivalent in every aspect, the adopted one offers more clarity in terms of code readability.

At this point we have all the necessary tools to create and run an `Observable`.

```
18 obs = newObservable $ \observer ->
19     do observer (OnNext 1)
20         observer (OnNext 2)
21         observer OnCompleted
22
23 main :: IO ()
24 main = subscribe obs print
25
26 {-
27 output>
28     OnNext 1
29     OnNext 1
30     OnCompleted
31 -}
```

Notice how, being a continuation, an `Observable` only pushes values once subscribed to and acts as a suspended computation otherwise.

The code above, being a toy example, fails to show some fundamental properties associated with this new interface; in particular, it fails to show how `Observable`s can actually handle asynchronous sources of data. The following snippet of code contains a more realistic and meaningful example of an `Observable` handling keyboard presses, asynchronous events by nature: whenever the user presses a key, an event containing the corresponding character is propagated to the `Observer` and will eventually be printed on the command line. It is worth noticing how our basic implementation of `Observable` based on continuations works just as well as a full blown one in terms of its core capability of handling asynchronous data.

```
17 obs :: Observable Char
18 obs = newObservable $ \observer -> do
19     keyboardCallback $= Just (\c p -> observer (OnNext c))
20
21 display :: DisplayCallback
22 display = do
23     clear [ ColorBuffer ]
24     flush
25
26 main :: IO ()
27 main = do
28     (_progName, _args) <- getArgsAndInitialize
29     _window <- createWindow "Observable Keyboard"
30     subscribe obs (print . show)
31     displayCallback $= display
32     mainLoop
```

This example brings us to the following observation: `Observable`s are capable of handling asynchronous data sources, yet the means by which the data is handled are not asynchronous by default. This is a common misconception and source of much confusion among the community: the `Observable` interface is not opinionated w.r.t. concurrency and therefore, by default, synchronously handles its incoming data, blocking the next incoming events whilst processing the current one. This behavior is not fixed though: as we will see in Chapter 3, it is possible to make use of `Scheduler`s to orthogonally introduce concurrency in our reactive systems, altering the control flow of the data processing allowing the user to dispatch the work to other threads.

At this point in the discussion we have arrived to a working implementation of a push based collection purely derived from mathematical and categorical concepts such as duality and continuations. In spite of being very insightful for theoretical discussions on the properties and relations of `Observable` and the continuation monad, this implementation of the reactive types is impractical in the context of a full fledged API. In the next Chapter we will take the necessary steps to build the bridge between theory and practice, providing a reference implementation of `Observable` more adapt to be utilized in real world applications.

For a full implementation of a Reactive Library based on the ideas presented in this section, aimed at highlighting the strong connection between `Observable` and other already existing functional structures from which it composes, see Appendice B.

OUT OF THE RABBIT HOLE: *Towards a usable API*

TBD

— Lewis Carol,
Alice in Wonderland

So far we focused our analysis on the essence of the `Observable` interface, setting aside the many operational concerns that would come up when trying to implement these concepts into a usable, commercial API. In this chapter we are going to build the bridge between our theoretical definition of `Observable` and a concrete and usable implementation of a reactive library, to which we will refer to as Rx.

In the remained of the discussion, we are going to introduce the *Reactive Contract*, a set of assumptions on the reactive types our library is going to build upon, *Schedulers*, which will allow us to bring concurrency into our reactive equation, *Subscriptions*, used to implement a mechanism for premature stream cancellation and, finally, *Operators*, the means that will make our reactive streams composable.

For the sake of clarity and completeness, the following set of interfaces represents the starting point for our discussion:¹

¹ `subscribe` has been renamed to `_subscribe` in order to both avoid naming conflicts later on in the discussion and reflect the fact that it should not be used from the outside.

```
1 newtype Observable a = Observable
2   { _subscribe :: Observer a -> IO ()
3   }
4 data Observer a = Observer
5   { onNext      :: a -> IO ()
6   , onError     :: SomeException -> IO ()
7   , onCompleted :: IO ()
8   }
```

It is worth noting that even though the Observable's theoretical foundations lie in the realm of functional programming, the road to make it usable is full of obstacles that are often better tackled using imperative programming features such as state. As much as I personally prefer a functional and pure approach to programming, I will favor, in the rest of the discussion, the solution that most clearly and easily solves the problem, be that functional or imperative.

3.1 The Reactive Contract

- Remember best effort cancellation.

3.2 Concurrency with Schedulers

At the end of Section 2.5 we discussed how `Observable` s, by default, handle data by means of a synchronous pipeline, blocking the processing of successive elements via the call stack. It is worth mentioning again how this synchronous processing does not affect the ability of `Observable` s to handle asynchronous data.

However, this synchronous behavior might not always be the best solution, especially in real world applications, where we might want to have a thread dedicated to listening to incoming events and one which processes them. Enter the `Scheduler` interface, an orthogonal^[?] structure w.r.t. `Observable` which allows us to introduce concurrency into our reactive equation.

`Scheduler` s allow us to alter the control flow of the data processing within an observable expression, introducing a way to dispatch the work of any number of operators to be executed within the specified context, e.g. a new thread.

The `Scheduler` interface looks like the following ².

²The interface presented in this section is the result of a simplification of the actual one, which involves `Subscription` s. We will discuss the impact of `Subscription` s on `Schedulers` in the next section; suffices to know


```
31 data Scheduler = Scheduler
32   { _schedule      :: IO () -> IO ()
33   , _scheduleDelay :: TimeInterval -> IO () -> IO ()
34   }
```

`Scheduler` s expose two functions which are essentially equal, modulo arbitrary delays in time. Both of these functions take an `IO` action as input and dispatch it to the appropriate execution context, producing a side effect.

To better understand `Scheduler` s, let us present the implementation of one of them, the `newThread` scheduler, which allows us to dispatch action to a new, dedicated thread.

```
36 newThread :: IO Scheduler
37 newThread = do
38   ch <- newTChanIO
39   tid <- forkIO $ forever $ do
40     join $ atomically $ readTChan ch
41     yield
42   return $ Scheduler (schedule_ ch) (scheduleDelay_ ch)
43   where
44     schedule_ ch io = atomically $ writeTChan ch io
45     scheduleDelay_ ch delay io = do
46       threadDelay $ toMicroSeconds delay
47       atomically $ writeTChan ch io
```

The `newThread` function gives us a side effectful way of creating a `Scheduler` by generating a new execution context - i.e. a new thread - and setting up the necessary tools for safe communication with it. The `Scheduler` functions we are provided, on the other hand, simply write the input `IO` action to the channel and return, effectively dispatching the execution of those actions to the new thread.

Up to this point we haven't mentioned `Observable` s at all. This is the reason why we previously claimed that `Scheduler` and `Observable` are connected by an orthogonal relationship: the two in-

that the version presented here has no negative effects w.r.t the generality of our discussion.

terfaces are independent from one another, yet, when used together within an observable expression, they provide the user with greater expressive power w.r.t. concurrency.

The only thing missing now is a way for us to combine the functionality of these two interfaces: `observeOn` and `subscribeOn` are the operators that will aid us on this task. The former will allow us to dispatch any call to an observer continuation on to the specified execution context, whereas the latter will allow us to control the concurrency of the `Observable` subscribe function.

For the sake of completeness and understandability, the following snippet contains a simple implementation of the `observeOn` operator together with a sample usage.

```
49 observeOn :: Observable a -> IO Scheduler -> Observable a
50 observeOn o sched = Observable $ \obr -> do
51   s <- sched
52   subscribe o (f s obr)
53   where
54     f s downstream = Observer
55       { onNext      = void . _schedule s . onNext downstream
56       , onError     = void . _schedule s . onError downstream
57       , onCompleted = void . _schedule s $ onCompleted downstream
58       }
```

```
60 obs = Observable $ \obr ->
61   do onNext obr 1
62     onNext obr 2
63     onNext obr 3
64     onCompleted obr
65
66 obr :: Observer Int
67 obr = Observer on oe oc
68   where
69     on v = do
70       tid <- myThreadId
71       print (show tid ++ ": " ++ show v)
72       oe = print . show
73       oc = print "Completed"
74
75 main :: IO ()
76 main = do
77   hSetBuffering stdout LineBuffering
78   subscribe obs' obr
79   tid <- myThreadId
80   putStrLn $ "MainThreadId: " ++ show tid
81   where
82     obs' = obs 'rxmap' (+1) 'observeOn' newThread 'rxmap' (+10)
83     rxmap = flip fmap
84
85 {-
86 output>
87   ThreadId 2: 12
88   ThreadId 2: 13
89   ThreadId 2: 14
90   Completed
91   MainThreadId: 1
92 -}
```

3.2.1 A note on the concept of time

Our discussion on push-based collections so far has not once mentioned the concept of time. This might appear strange, especially to the reader familiar with Functional Reactive Programming, where functions over continuous time are at the foundations of the theory. As we mentioned in Section 1.3.3, this dependency on continuous time comes at a great cost: commercial FRP libraries fail to successfully implement the concepts found in the theory as they cannot avoid simulating continuous time and approximating functions operating over it, being this concept inherently discrete in the context of computers.

Rx, on the other hand, completely sheds the notion of time from the notion of reactivity^[?] , shifting its focus, with the help of `Scheduler s`, to concurrency instead. Time still plays a role, although indirect, within the library: events are processed in the order they happen, and operators make such order is maintained, ultimately handing over to the user a stream of time-ordered events. See Section 3.1 for a broader discussion on this matter.

3.2.2 Impact on Operators

Previously we discussed how introducing concurrency is an orthogonal concept w.r.t. our reactive library. This statement is only partly true and falls short of its promises when we are dealing with combiner operators (see Section 3.4) such as `(>=)` or `combineLatest`. These operators will not work at their full potential in a synchronous setting, due to the fact that subscribing to a stream will consume it entirely - or forever process, in the case of an infinite stream - before allowing the operator to subscribe to a different one, effectively making interleaving of events impossible.

The problem is gracefully solved with the introduction of `Scheduler s`, which, by allowing for `Observable s` to be executed on different contexts, indirectly make it possible for interleaving to happen and for combiner operators to work at their full potential. This comes at a - even though small - cost, which motivates our claim of partial orthogonality: the internal state of a combiner operator needs to be synchronized and carefully handled as, with the introduction of concurrency, it can now be modified from different execution contexts.

3.3 Subscriptions

- now purely practical concern: stop stream while processing from anywhere, tell observable I dont want to observe anymore, i.e. the observer is done. - what is it: IO action/s | in theory can be seen as `readerT`, the subscription is always the same one, but you can add stuff to it. - who does it belong to: the observer, you get one for each subscribe. - how does it work: call `unsubscribe`, best effort stop work - impact on operators: a need for children -> now you can free resources - impact on schedulers:

kill the thread - write that it only works now that we have schedulers - edge enforcement - subject + takeuntil - why cannot be monoid, we need remove

With schedulers, we are now able to handle observable streams from different execution contexts. The next step is to add a mechanism that will allow us to stop a stream from anywhere in our program, whenever we don't require it's data - i.e. a mechanism to tell the `Observable` that one of it's `Observer`s is not interested in receiving events anymore. We are going to do so by augmenting our informal definition of `Observable` from section 2.5 with a new function `(%)`, which incorporates the notion of a cancellation mechanism through the `Subscription` datatype:

```
1 %a  :: a -> IO Subscription
2 %?a :: (Either SomeException (Maybe a) -> IO ()) -> IO Subscription
```

With this change, each execution of the `Observable` function now returns a `Subscription`, a means for the user to prematurely terminate the processing of the stream.

We discussed in the previous section how schedulers effectively boost the expressive power of our reactive expressions by introducing concurrency and interleaving among events originating from different streams. Introducing a cancellation mechanism, on the other hand, is a purely practical concern: although very useful from a practical perspective, especially in the context of resource management, it doesn't impact expressive power from a reactive point of view.

The first step in designing this new feature is to understand how the already existing interfaces will be affected by the newly introduced one; we just saw how the `_subscribe` function should now return an action resulting in an instance of type `Subscription`. The next question is the following: to whom does a `Subscription` belongs to? The key observation in addressing this question is that an `Observable` can be subscribed to by multiple `Observer`s; our goal is to provide a mechanism that will allow for a fine-grained control over which `Observer` is supposed to stop receiving events. The answer is then straightforward: the notion of subscription is tight to that of observer. The following snippet incorporates this notion:

```
1 $a  :: (Subscription, Either SomeException (Maybe a) -> IO ())
2 %$a :: (Subscription, Either SomeException (Maybe a) -> IO ())
3     -> IO Subscription
```

Let's quickly summarize what we have discussed so far: a subscription is some object which will allow us to prematurely stop observing a specific stream; since any stream can be subscribed to by multiple observers, we need to associate subscriptions to observers as opposed to observables. Lastly, a subscription is returned every time an observer is subscribed to a stream through the `_subscribe` function. The following modifications to our reactive interfaces reflect these ideas:

```
12 newtype Observable a = Observable
13     { _subscribe :: Observer a -> IO Subscription
14     }
15
16 data Observer a = Observer
17     { onNext      :: a -> IO ()
18     , onError     :: SomeException -> IO ()
19     , onCompleted :: IO ()
20     , subscription :: Subscription
21     }
```

do not subscribe an observer to 2 different streams

So far we have talked a lot about `Subscription`s, yet we haven't clarified what the type really looks like. The general idea is to have `Subscription` record the state of the `Observer` w.r.t. the `Observable` it is subscribed to - be that subscribed or unsubscribed. This can be easily achieved with a variable `_isUnsubscribed :: IORef Bool` initialized to `False`, indicating that the associated `Observer` is initially not unsubscribed.

From a practical point of view, it is useful to augment `Subscription` with some additional functionality. The following code shows a definition of `Subscription` which incorporates an `IO ()` action to be executed at unsubscription time. This is particularly useful when we want to associate resource cleanup actions to the termination - be that forced or natural - of a stream observation.

```
1 data Subscription = Subscription
2     { _isUnsubscribed :: IORef Bool
3     , onUnsubscribe   :: IO ()
4     }
```

We will see, later on, how we can further augment this interface in order to incorporate additional functionality useful for internal operator coordination as well as management of schedulers.

It's now time to introduce the two functions at the hearth of the whole cancellation mechanism: `unsubscribe` will take care of modifying the state carried by the `Subscription` - i.e. setting `_isUnsubscribed` to `False` - whereas `subscribe` will act as a proxy for the `_subscribe` function from `Observable`, calling it with a special `Observer` which will prevent events from propagating to the downstream once the `Subscription` has been unsubscribed. Note how, for this mechanism to work, it is necessary the two `Observer`s to share the same `Subscription`.

```

71 subscribe :: Observable a -> Observer a -> IO Subscription
72 subscribe obs obr = do
73     _subscribe obs $ Observer on' oe' oc' s
74     where
75         s      = subscription obr
76         oc'    = ifSubscribed $ onCompleted obr >> unsubscribe s
77         on' a  = ifSubscribed $ onNext obr a
78         oe' e  = ifSubscribed $ finally (onError obr e) (unsubscribe s)
79         ifSubscribed a = isUnsubscribed s >=> flip unless a
80
81 unsubscribe :: Subscription -> IO ()
82 unsubscribe s = do
83     writeIORef (_isUnsubscribed s) True
84     onUnsubscribe s

```

So, what is a `Subscription` then? Conceptually, it could be a simple boolean recording whether an `Observer` is unsubscribed from the stream. In practice it is useful to augment it with some additional information. First of all, `Subscription`s can be used for resource cleanup at the termination - either forced or natural - of a stream. It is therefore useful to associate an `IO ()` action to the datatype, which will represent the computation to be executed once the stream is unsubscribed from. Another improvement is making the type recursive, i.e. allowing `Subscription`s to contain, add and remove children subscriptions. This will be extremely useful for internal coordination of operators such as `(>=)` and for resource management in the context of schedulers - more on this in the next sub-subsections.

We could just chain IO actions, but then we cant remove, which can lead to stack overflow with huge thunks in flatmap.

Talk about the best effort in canceling work and eventual consistency with the contract.

```
34 data Subscription = Subscription
35   { onUnsubscribe    :: IO ()
36   , _isUnsubscribed  :: IORef Bool
37   , subscriptions    :: IORef [Subscription]
38   }
```

As a design decision, when unsubscribe is called on an observable subscription, the observable sequence will make a best effort attempt to stop all outstanding work. This means that any queued work that has not been started will not start. Any work that is already in progress might still complete as it is not always safe to abort work that is in progress. Results from this work will not be signaled to any previously subscribed observer instances.

Motivation for children subscriptions: some operators will create inner observables and therefore will need to unsubscribe from them when the outer observable is unsubscribed from.

On why using `Cont () (StateT Subscription IO) (Event a)` wouldn't work: it would be perfect in order to thread Subscriptions throughout execution making it usable inside operators, since the call to the continuation would return a state which would then be sequenced by `»=`. This method fails with schedulers, in particular `newThread` since the state of the action executed on the new thread would be disconnected from the threading mentioned above. The other thread would get a state but it would not know what to do with it and would not have any ways to connect it to the original one passed by the subscribe. Another way is to use `mapStateT` to map the IO action that will result from the state to an action on the other thread. Again, this method won't work,

On the motivation for a Subscription: it is needed so the user can cancel work at any time. This implies that a scheduler is used. If this is not the case the subscription will be returned synchronously after the execution of the whole stream. It will therefore be already unsubscribed and calling unsubscribe will be a NoOp. In the case that we actually use schedulers then we can unsubscribe from anywhere in our program.

Now the question is the following: can we reproduce the behaviour of unsubscribe with operators so that we can eliminate subscriptions altogether? That is, we can hide them to the outside and use them only inside the stream, we still need them but we don't necessarily need to return them when we subscribe. The behaviour can be easily replaced by the use of `takeUntil(Observable a)` where we pass in a subject that will be signaled when we want to stop the stream. The `takeUntil` operator fires events from the upstream up until the point in which we signal the subject, then stops the stream and unsubscribes.

With this approach we seemingly lose one thing, i.e. the ability to specify what happens at unsubscribing time; this functionality can simply be regained by a subscribe function that takes the unsubscribe action as a parameter and runs it when the stream is unsubscribed.

3.4 Operators

The first thing we want to re-introduce is handling exceptions and termination of a stream. Where an `Iterable` can return a value, terminate or throw an exception when we ask for a value, and `Observable`, being its dual, can produce one or more values, terminate or throw an exception when it is subscribed to.

A more appropriate type for our interface is then the following.

```
4 type Observer a = Either SomeException (Maybe a) -> IO ()
5 type Observable a = ContT () IO (Observer a)
```

Remove this Event shit and move to the custom type directly? Not sure, I like the discussion on bind.

Now, this code is not exactly the definition of readable; let's apply some good design skills to make it more pleasant to the eye without changing its meaning.

```
4 -- datatype representing Either SomeException (Maybe a)
5 data Event a =
6     OnNext a
7   | OnError SomeException
8   | OnCompleted
9
10 type Observer a = Event a -> IO ()
11 type Observable a = ContT () IO (Event a)
```

Although this might not look like a big change, it greatly influences the design of our API. We are, in fact, changing our instantiation of the continuation monad to an input type that is not `a` anymore, but `Event a`. On the other hand, our type variable for `Observable` is still `a`. This is not an issue

per se, but it has one big consequence: the flatmap function that we inherit from the continuation monad is not the one that we want to expose from our API anymore. The types differ like so.

```
13  -- bind inherited from the Continuation monad
14  (>>=) :: Observable a -> (Event a -> Observable b) -> Observable b
15
16  -- bind that we would like to expose from our API
17  flatmap :: Observable a -> (a -> Observable b) -> Observable b
```

This has many implications, first of all, we are gonna need to implement flatmap by ourselves.. see todo below...

It has now come the time to move away from an implementation of Observable as a type synonym. We have already seen how the current implementation using `Event a` does not allow for a correspondence between `>>=` operations; this will only create confusion in the future. The next step is then to define our own observable type, which will clearly be really similar to the Continuation monad and subsequently prove that it is itself a monad.

```
newtype Observable a = Observable { subscribe :: Observer a -> IO () }
data Observer a = Observer
  { onNext      :: a -> IO ()
  , onError     :: SomeException -> IO ()
  , onCompleted :: IO ()
  }
}
```

With this implementation we have eliminated the materialisation of the event types. The Observer is now not a single function from `Event a -> IO ()` but a collection of 3 continuations that will be used inside the observable depending on the type of the event. It is clear that this implementation of Observable has not changed in functionality from the previous one using the Continuation Monad, it has just dematerialized the 3 types of events in 3 functions which handle them.

The next step is to make Observable a monad

```
instance Monad Observable where
    return a = observable (\obr -> onNext obr a)
    o >>= f = ...
```

The return function is the exact same as in the continuation monad, with the only difference that we have now 3 continuations to chose from instead of a single one.

Bind, on the other hand, is completely different from the Cont monad implementation; in this case ...

finish the discussion

The only thing left to do now is to prove the monad laws to show that Observable really is a monad.

```
19 -- Monad Laws
20 -- return a >>= k = k a
21 -- m >>= return = m
22 -- m >>= (x -> k x >>= h) = (m >>= k) >>= h
```

We mentioned before how the bind from Cont differs from our in the Observable. Below I will show that in this implementation it corresponds to a function lift that ...

Talk about lift = »= in Cont.

By using lift we can transform streams and implement operators...

Modify keyboard press example from before to handle errors and termination. Point to later discussion regarding the rx contract, since now we can detect termination and errors but there is no guarantee that nothing will come after we receive them, i.e. that we abide the contract.

TODO LIST

Write overview	2
Github link	2
Probably put the derivation of Reactive Streams in the appendix, it's not really the focus, I guess.	7
do not subscribe an observer to 2 different streams	30
We could just chain IO actions, but then we cant remove, which can lead to stack overflow with huge thunks in flatmap.	31
Talk about the best effort in canceling work and eventual consistency with the contract. . .	31
Motivation for children subscriptions: some operators will create inner observables and therefore will need to unsubscribe from them when the outer observable is unsubscribed from.	32
Remove this Event shit and move to the custom type directly? Not sure, I like the discussion on bind.	33
finish the discussion	35
Talk about lift = »= in Cont.	35
Modify keyboard press example from before to handle errors and termination. Point to later discussion regarding the rx contract, since now we can detect termination and errors but there is no guarantee that nothing will come after we receive them, i.e. that we abide the contract.	35



Begins an appendix

A.1 Categorical Duality

A.2 Continuations

A.3 (Co)Products

A.4 (Un)Currying

A.5 (Covariance & Contravariance

A.6 Functors



APPENDIX B

Begins an appendix
Formal RX Implementation

BIBLIOGRAPHY

- [1] *Scopus search: Reactive programming 1960-2016.*
- [2] E. BAINOMUGISHA, A. L. CARRETON, T. V. CUTSEM, S. MOSTINCKX, AND W. D. MEUTER, *A survey on reactive programming*, ACM Computing Surveys (CSUR), 45 (2013), p. 52.
- [3] A. BENVENISTE AND G. BERRY, *The synchronous approach to reactive and real-time systems*, Proceedings of the IEEE, 79 (1991), pp. 1270–1282.
- [4] J. EDWARDS, *Coherent reaction*, in Proceedings of the 24th ACM SIGPLAN conference companion on Object oriented programming systems languages and applications, ACM, 2009, pp. 925–932.
- [5] D. FAHLAND, D. LÜBKE, J. MENDLING, H. REIJERS, B. WEBER, M. WEIDLICH, AND S. ZUGAL, *Declarative versus imperative process modeling languages: The issue of understandability*, in Enterprise, Business-Process and Information Systems Modeling, Springer, 2009, pp. 353–366.
- [6] B. FURHT AND A. ESCALANTE, *Handbook of cloud computing*, vol. 3, Springer, 2010.
- [7] E. GAMMA, *Design patterns: elements of reusable object-oriented software*, Pearson Education India, 1995.
- [8] T. GOWERS, J. BARROW-GREEN, AND I. LEADER, *The Princeton companion to mathematics*, Princeton University Press, 2010.
- [9] C. H. GROUP, *Covariance, contravariance, and positive and negative position.*
- [10] D. LANEY, *3d data management: Controlling data volume, velocity and variety*, META Group Research Note, 6 (2001), p. 70.
- [11] E. MEIJER, *Your mouse is a database*, Queue, 10 (2012), p. 20.
- [12] E. MEIJER, *Reactive streams keynote*. LambdaJam, 2014.

BIBLIOGRAPHY

- [13] E. MEIJER, K. MILLIKIN, AND G. BRACHA, *Spicing up dart with side effects*, Queue, 13 (2015), p. 40.
- [14] J. NEWBERN, *All about monads*.
- [15] TEL, *What's up with contravariant?*